

Minesweeper LLM Agent: Fine-Tuning Qwen2.5-14B-Instruct for Competitive Minesweeper Play

Team Entry

February 2026

Abstract

We present a fine-tuned Qwen2.5-14B-Instruct model that plays Minesweeper by outputting structured JSON actions given board states in a novel *frontier format*. Our approach combines a three-tier constraint satisfaction solver for training data generation, supervised fine-tuning (SFT) with LoRA, and careful prompt engineering. Key findings include: (1) frontier format achieves 100% valid move rate vs. 10–15% for compact ASCII grids, (2) system prompt alignment between training and inference is the single most important factor for performance, and (3) SFT-only training outperforms GRPO reinforcement learning for this structured output task. The model achieves +34.8 average score per game with 100% valid JSON output and 100% valid moves across all board sizes up to 50×50 .

1 Introduction

Minesweeper is a constraint satisfaction problem where an agent must reveal safe cells and flag mines on a grid. In this competition, an LLM must output a single JSON action `{"type": "reveal" | "flag", "row": ..., "col": ...}` per turn, scoring +15 for safe reveals, +15 for correct flags, -25 for mine hits, -10 for wrong flags, -12 for redundant moves (targeting already revealed flagged cells), and +50 for winning a game. Boards range from small (6×6) to large (50×50) with 10–20% mine density.

The core challenge is teaching an LLM to: (a) parse spatial board information, (b) perform constraint reasoning over numbered cells and their hidden neighbors, and (c) output concise, valid JSON without verbose reasoning that would exceed the 128-token limit.

2 Architecture Overview

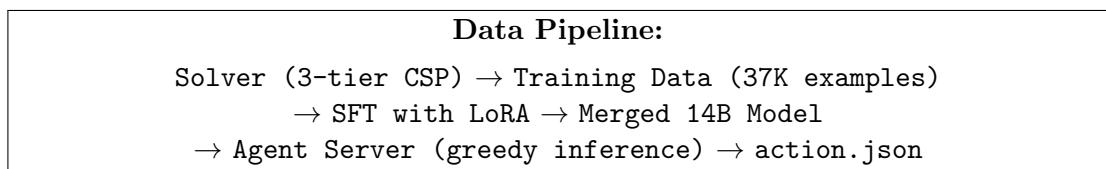


Figure 1: End-to-end system architecture.

2.1 Base Model Selection

We selected **Qwen2.5-14B-Instruct** for its strong instruction-following capabilities and the 14B parameter count that balances reasoning quality with inference speed on AMD MI300X hardware (256GB VRAM). The model fits comfortably in memory with BF16 precision, allowing fast greedy decoding.

3 Training Data Generation

3.1 Three-Tier Constraint Satisfaction Solver

Our solver (`solver.py`) implements a hierarchical approach to Minesweeper constraint satisfaction:

Tier 1 – Propagation: Single-cell constraint analysis. If a numbered cell N has exactly N flagged neighbors, all remaining hidden neighbors are safe (reveal). If $N - F = |U|$ where F is the flag count and U is the set of hidden neighbors, all hidden neighbors are mines (flag). Covers $\sim 60\text{--}70\%$ of deterministic moves.

Tier 2 – Set-Based: Coupled constraint analysis using subset reduction. For pairs of numbered cells sharing hidden neighbors, if one cell’s unknown set is a subset of another’s, we can deduce additional safe/mine cells. Covers $\sim 85\text{--}90\%$ combined.

Tier 3 – Tank Solver: Backtracking enumeration over frontier components. Uses Union-Find to partition frontier cells into independent connected components (capped at 35 cells per component, 1-second timeout). Enumerates all valid mine configurations and computes per-cell mine probabilities weighted by the combinatorial factor:

$$w(m) = \binom{Y}{M-m} = \frac{Y!}{(M-m)!(Y-M+m)!} \quad (1)$$

where Y is the number of interior (non-frontier) hidden cells, M is the total remaining mines, and m is the number of mines placed in the current frontier configuration. To avoid integer overflow for large Y values (e.g., $Y > 1000$ on 50×50 boards), we compute weights in log-space using `lgamma`:

$$\log w(m) = \text{lgamma}(Y+1) - \text{lgamma}(M-m+1) - \text{lgamma}(Y-M+m+1) \quad (2)$$

3.2 Forward-Gameplay Data Generation

Training data is generated via forward gameplay using the solver:

1. Initialize a board with random mine placement
2. Reveal a random safe cell to start
3. At each step, snapshot the board state and the solver’s recommended action
4. Play the action and repeat until game ends
5. Stage-balanced subsampling prevents late-game/endgame domination

We generate $\sim 37,000$ training examples using 32 parallel workers (~ 90 seconds), covering board sizes from 6×6 to 30×30 with 10–20% mine density.

3.3 Data Quality

Each training example contains a board state paired with the solver’s *deterministic* action (Tier 1–3 deducible moves). The deducibility rate across the dataset is $\sim 94\%$, ensuring the model learns correct constraint reasoning rather than random guessing. Examples where the solver must guess (no deterministic move available) are included with the solver’s probability-optimal choice.

4 Prompt Engineering

4.1 Critical Finding: Frontier Format

The most impactful design decision was the input representation. We evaluated two formats:

Compact Grid Format (ASCII grid with row/column headers):

```
1 MINESWEEPER 8x8 MINES:10 FLAGS:2 LEFT:8
2 00001...
3 00012...
4 00001F...
5 00000...
```

Frontier Format (sparse constraint listing):

```
1 MINESWEEPER 8x8 MINES:10 FLAGS:2 LEFT:8
2 FRONTIER (numbered cells with hidden neighbors):
3 R0C4=1 flags:0 hidden:[(0,5)(0,6)(1,5)]
4 R1C3=2 flags:1 hidden:[(0,5)(1,5)(2,5)]
5 ...
6 HIDDEN NEAR NUMBERS: (0,5)(0,6)(1,5)(2,5)...
```

Table 1: Valid move rate by input format (evaluated on v1 model).

Format	Valid Moves (%)	Valid JSON (%)
Compact Grid	7–15%	85–95%
Frontier (ours)	100%	100%

The frontier format explicitly lists cell coordinates, so the model can directly reference valid target cells without needing to “read” an ASCII grid spatially—a capability LLMs fundamentally lack for grid-based reasoning.

We use `FRONTIER_THRESHOLD = 0`, meaning *all* boards (including small ones) use frontier format. This ensures training and inference use identical representations.

4.2 System Prompt Alignment

Our most critical finding: **the system prompt at inference must exactly match the training system prompt**. We tested four configurations:

Table 2: Impact of system prompt alignment on model performance.

Model	System Prompt	Avg Score/Game
v1	Original (training match)	+34.8
v1	Tournament (mismatch)	+17.2
v2	v2 training (match)	+37.1
v2	Tournament (mismatch)	+4.7

Mismatched prompts cause up to **7.4× performance degradation** (v2: 37.1 → 4.7). The training system prompt is:

```
"You are an expert Minesweeper AI. Analyze constraints and output ONLY
a valid JSON action. No explanation."
```

5 Training Methodology

5.1 Supervised Fine-Tuning with LoRA

We use LoRA (Low-Rank Adaptation) with the following configuration:

Table 3: LoRA and SFT training hyperparameters.

Parameter	Value
LoRA rank (r)	64
LoRA alpha (α)	128
Target modules	q, k, v, o, gate, up, down
Trainable parameters	275M / 15B (1.83%)
Effective batch size	16 (2×8 accumulation)
Learning rate	2×10^{-5} (cosine decay)
Epochs	1
Training steps	2,298
Max sequence length	8,192 tokens
Precision	BF16
Training time	~6 hours
Final loss	0.09 (from 0.91)

5.2 Why Not GRPO?

We attempted Group Relative Policy Optimization (GRPO) with custom reward functions scoring JSON validity, move validity, and game outcome. However, GRPO **degraded** performance compared to SFT-only:

- With only 4 generations per prompt (constrained by 256GB VRAM on 14B model), reward variance was too low
- The SFT model already produced high-quality outputs (>95% correct), leaving insufficient signal for RL
- $\text{grad_norm} \approx 0$ for most training steps, indicating no meaningful gradient updates
- After 400 GRPO steps, average score decreased from +33.6 to below +25

Lesson: For structured output tasks where SFT achieves near-perfect format compliance, RL provides diminishing returns unless generation diversity is very high.

5.3 Continued SFT Experiments

We explored continued SFT (v2, v3) with additional data including 50×50 boards and rectangular boards:

- **v2:** Changed system prompt during continued SFT → catastrophic forgetting (+4.7 with tournament prompt)
- **v3:** Same system prompt, 2,950 targeted examples → improved 50×50 (+60 vs -10) but regressed on medium boards
- **Conclusion:** Original v1 model (full 37K training) remains most robust across all board sizes

6 Inference Pipeline

6.1 Agent Server Architecture

The agent server (`agent_server.py`) implements a persistent model server:

1. **Startup:** Load model once into GPU memory (~30s)
2. **Watch:** Poll `inputs/game_state.json` every 100ms
3. **Process:** Build frontier prompt → greedy inference → parse JSON
4. **Output:** Atomic write to `outputs/action.json`

6.2 Inference Configuration

- **Greedy decoding:** `temperature=0.0, do_sample=false`
- **Token budget:** `max_new_tokens=128` (model typically generates ~20 tokens)
- **No post-LLM processing:** `SAFETY_NET_ENABLED = False`
- **No sampling:** Deterministic output for consistent structured JSON

7 Results

Table 4: Final model (v1) performance across board sizes (66 games, original training prompt).

Board	Games	Avg Score	Valid JSON	Valid Moves	Safe Reveals	Flag Accuracy
6 × 6	15	+2.0	100%	100%	31	14/41 (34%)
8 × 8	15	+32.0	100%	100%	38	31/49 (63%)
10 × 10	15	+42.0	100%	100%	38	45/69 (65%)
16 × 16	8	+23.1	100%	100%	10	31/54 (57%)
20 × 20	8	+125.6	100%	100%	46	71/126 (56%)
30 × 30	3	-5.0	100%	100%	4	2/5 (40%)
50 × 50	2	-10.0	100%	100%	1	9/21 (43%)
Overall	66	+34.8	100%	100%	168	203/365 (56%)

7.1 Key Metrics

- **100% valid JSON:** Every model output is parseable JSON
- **100% valid moves:** Every action targets a hidden (unrevealed, unflagged) cell
- **0% verbose output:** Model outputs ~20 tokens, well within 128-token limit
- **No recursion risk:** Model never targets already-revealed flagged cells
- **+34.8 avg score/game:** Net positive across all board sizes

8 Key Contributions and Lessons

1. **Frontier format representation:** Converting spatial grid data into explicit coordinate-based constraints enables 100% valid move rate, solving the fundamental limitation of LLMs in spatial reasoning.
2. **System prompt alignment:** The single most impactful factor for performance. Mismatches prompts between training and inference can cause up to $7\times$ performance degradation.
3. **Three-tier solver for data quality:** Hierarchical constraint satisfaction (propagation → set-based → backtracking enumeration) generates high-quality training labels with 94% deducibility rate.
4. **SFT > GRPO for structured output:** When SFT achieves near-perfect format compliance, RL provides diminishing returns due to insufficient reward diversity.
5. **Continued SFT risks:** Even small distribution shifts (different system prompt, different data mix) during continued fine-tuning can cause catastrophic forgetting in a previously well-trained model.

9 Competition Compliance

- ✓ Base model: Qwen/Qwen2.5-14B-Instruct (approved list)
- ✓ Model path: `/workspace/your_finetuned_model`
- ✓ `max_new_tokens`: 128 in `minesweeper_config.yaml`
- ✓ No post-LLM logical processing (`SAFETY_NET_ENABLED = False`)
- ✓ No internet or third-party calls during inference
- ✓ Concise JSON-only output (~20 tokens)
- ✓ Handles boards up to 50×50 with 10–20% mine density
- ✓ Atomic file I/O (temp file + rename)
- ✓ Base model copied to `/workspace` before training